

Project_1_regression

Muhammad Zubair

10/14/2021

Link to the dataset: <https://www.kaggle.com/zaynshahbaz/pakistan-car-prices>

```
# Reading in the data
df <- read.csv("updated_pakwheels.csv")
```

Data Exploration

```
# Viewing the first 5 rows of dataset
head(df)
```

```
##      Ad.No      Name      Price Model.Year
## 1 4096758 Toyota Vitz F 1.0 2017 2385000      2017
## 2 4168305 Toyota Corolla GLi Automatic 1.3 VVTi 2019 111000      2019
## 3 4168298 Suzuki Alto VXL 2019 1530000      2019
## 4 4168307 Suzuki Alto VXR 2019 1650000      2019
## 5 4168306 Toyota Corolla XLi VVTi 2010 1435000      2010
## 6 4168303 Honda Civic 1.5 RS Turbo 2017 3850000      2017
##      Location Mileage Registered.City Engine.Type
## 1      G- 8, Islamabad Islamabad      9869 Un-Registered      1
## 2      Peshawar KPK      11111 Islamabad      1
## 3      Akora Khattak, Nowshera KPK      17500 Un-Registered      1
## 4      Abdullahpur, Faisalabad Punjab      9600 Lahore      1
## 5      9th Avenue, Islamabad Islamabad      120000 Islamabad      1
## 6      Peshawar Road, Rawalpindi Punjab      22000 Islamabad      1
##      Engine.Capacity Transmission Color Assembly Body.Type
## 1      1000      1 Silver Imported      1
## 2      1300      1 White Local      2
## 3      660      1 White Local      1
## 4      660      2 White Local      1
## 5      1300      2 Black Local      2
## 6      1500      1 Black Local      2
##
## 1      ABS, AM/FM Radio, Air Bags, Air Conditioning, CD Player, DV
## 2      ABS, AM/FM Radio, Air Bags, Air Conditioning, CD Player, DV
## 3      ABS, AM/FM Radio, Air Bags, Air Conditioning, CD Player, DV
## 4      ABS, AM/FM Radio, Air Bags, Air Conditioning, CD Player, DV
## 5      ABS, AM/FM Radio, Air Bags, Air Conditioning, CD Player, DV
```

```
## 6 ABS, AM/FM Radio, Air Bags, Air Conditioning, Alloy Rims, CD Player, Cruise Control, DVD Player,
## Last.Updated
## 1 11-Jul-20
## 2 12-Jul-20
## 3 12-Jul-20
## 4 12-Jul-20
## 5 12-Jul-20
## 6 12-Jul-20
##
## URL
## 1 https://www.pakwheels.com/used-cars/toyota-vitz-2017-for-sale-in-islamabad-4096758
## 2 https://www.pakwheels.com/used-cars/toyota-corolla-2019-for-sale-in-peshawar-4168305
## 3 https://www.pakwheels.com/used-cars/suzuki-alto-2019-for-sale-in-nowshera-4168298
## 4 https://www.pakwheels.com/used-cars/suzuki-alto-2019-for-sale-in-faisalabad-4168307
## 5 https://www.pakwheels.com/used-cars/toyota-corolla-2010-for-sale-in-islamabad-4168306
## 6 https://www.pakwheels.com/used-cars/honda-civic-2017-for-sale-in-rawalpindi-4168303
```

```
# Viewing the last 5 rows of data
tail(df)
```

```
## Ad.No Name Price Model Year
## 46018 3806954 Honda Civic Oriel 1.8 i-VTEC CVT 2017 3300000 2017
## 46019 3448128 Honda Vezel Hybrid X 2015 3400000 2015
## 46020 3737684 Toyota Aqua S 2015 2450000 2015
## 46021 3349017 Honda Civic VTi Prosmatec 1.8 i-VTEC 2015 3250000 2015
## 46022 3748215 Toyota Aqua G 2016 3000000 2016
## 46023 3806951 Toyota Corolla GLi Automatic 1.3 VVTi 2015 2250000 2015
## Location Mileage Registered.City Engine.Type Engine.Capacity
## 46018 Gujranwala Punjab 40000 Lahore 1 1800
## 46019 Lahore Punjab 32000 Un-Registered 1 1500
## 46020 Rawalpindi Punjab 52000 Un-Registered 1 1500
## 46021 Lahore Punjab 125000 Lahore 1 1800
## 46022 Gujranwala Punjab 60000 Lahore 1 1500
## 46023 Gujranwala Punjab 77000 Gujranwala 1 1300
## Transmission Color Assembly Body.Type Features Last.Updated
## 46018 1 Black Local 2 14-Jan-20
## 46019 1 Black Imported 5 28-Jul-19
## 46020 1 Blue Imported 1 18-Dec-19
## 46021 1 Black Local 2 4-Jun-19
## 46022 1 Black Imported 1 22-Dec-19
## 46023 1 Assembly Local 2 14-Jan-20
```

```
## URL
## 46018 https://www.pakwheels.com/used-cars/honda-civic-2017-for-sale-in-gujranwala-3806954
## 46019 https://www.pakwheels.com/used-cars/honda-vezel-2015-for-sale-in-lahore-3448128
## 46020 https://www.pakwheels.com/used-cars/toyota-aqua-2015-for-sale-in-rawalpindi-3737684
## 46021 https://www.pakwheels.com/used-cars/honda-civic-2015-for-sale-in-lahore-3349017
## 46022 https://www.pakwheels.com/used-cars/toyota-aqua-2016-for-sale-in-gujranwala-3748215
## 46023 https://www.pakwheels.com/used-cars/toyota-corolla-2015-for-sale-in-gujranwala-3806951
```

```
# Dimensions of our data
dim(df)
```

Our dataset have 46k rows and 16 attributes

```
## [1] 46023    16
```

```
# Data types for each column in our dataset
str(df)
```

Columns with long descriptions and sentences will need to be dropped, also the rest of the columns will be converted into factor variable

```
## 'data.frame':    46023 obs. of  16 variables:
## $ Ad.No          : int  4096758 4168305 4168298 4168307 4168306 4168303 4168304 4168309 4168310 4168311 ...
## $ Name           : chr  "Toyota Vitz F 1.0 2017" "Toyota Corolla GLi Automatic 1.3 VVTi 2019" "Suzuki Swift 1.5 2017" ...
## $ Price          : int  2385000 111000 1530000 1650000 1435000 3850000 1440000 1425000 2650000 3350000 ...
## $ Model.Year     : int  2017 2019 2019 2019 2010 2017 2017 2012 1998 2017 ...
## $ Location       : chr  "G- 8, Islamabad Islamabad" "Peshawar KPK" "Akora Khattak, Nowshera KPK" ...
## $ Mileage        : int  9869 11111 17500 9600 120000 22000 31000 101000 110000 60000 ...
## $ Registered.City: chr  "Un-Registered" "Islamabad" "Un-Registered" "Lahore" ...
## $ Engine.Type    : int  1 1 1 1 1 1 1 2 1 ...
## $ Engine.Capacity: int  1000 1300 660 660 1300 1500 1000 1000 3000 1800 ...
## $ Transmission   : int  1 1 1 2 2 1 2 1 1 1 ...
## $ Color          : chr  "Silver" "White" "White" "White" ...
## $ Assembly       : chr  "Imported" "Local" "Local" "Local" ...
## $ Body.Type      : int  1 2 1 1 2 2 1 1 3 2 ...
## $ Features       : chr  "ABS, AM/FM Radio, Air Bags, Air Conditioning, CD Player, DVD Player, Immobilizer" ...
## $ Last.Updated   : chr  "11-Jul-20" "12-Jul-20" "12-Jul-20" "12-Jul-20" ...
## $ URL            : chr  "https://www.pakwheels.com/used-cars/toyota-vitz-2017-for-sale-in-islamabad"
```

```
# Running some stats on the dataset
summary(df)
```

```
##      Ad.No      Name      Price      Model.Year
## Min.   : 13381  Length:46023  Min.   : 111000  Min.   :1990
## 1st Qu.:4051758  Class :character  1st Qu.: 850000  1st Qu.:2007
## Median :4103354  Mode  :character  Median : 1450000  Median :2013
## Mean   :4070389                Mean   : 2014144  Mean   :2011
## 3rd Qu.:4142396                3rd Qu.: 2300000  3rd Qu.:2016
## Max.   :4168339                Max.   :77500000  Max.   :2019
##      Location      Mileage  Registered.City  Engine.Type
## Length:46023      Min.   :    1  Length:46023  Min.   :1.000
## Class :character  1st Qu.: 48892  Class :character  1st Qu.:1.000
## Mode  :character  Median : 80000  Mode  :character  Median :1.000
##                      Mean   : 90964                Mean   :1.084
##                      3rd Qu.:120000                3rd Qu.:1.000
##                      Max.   :999999                Max.   :3.000
##      Engine.Capacity  Transmission      Color      Assembly
## Min.   : 16          Min.   :1.000  Length:46023  Length:46023
## 1st Qu.:1000         1st Qu.:1.000  Class :character  Class :character
## Median :1300         Median :2.000  Mode  :character  Mode  :character
## Mean   :1313         Mean   :1.535
```

```
## 3rd Qu.:1500    3rd Qu.:2.000
## Max.      :6600    Max.      :2.000
##   Body.Type      Features      Last.Updated      URL
## Min.      :1.000    Length:46023    Length:46023    Length:46023
## 1st Qu.:1.000    Class :character    Class :character    Class :character
## Median :2.000    Mode  :character    Mode  :character    Mode  :character
## Mean      :1.772
## 3rd Qu.:2.000
## Max.      :6.000
```

```
# Checking for null values in our data
sapply(df, function(x) sum(is.na(x)))
```

```
##           Ad.No           Name           Price           Model.Year           Location
##              0              0              0              0              0
##      Mileage Registered.City      Engine.Type Engine.Capacity      Transmission
##              0              0              0              0              0
##           Color           Assembly      Body.Type           Features      Last.Updated
##              0              0              0              0              0
##           URL
##              0
```

Data cleaning

Link to dataset: <https://www.kaggle.com/zaynshahbaz/pakistan-car-prices>

```
df <- subset(df, select = -c(Ad.No, Last.Updated, URL, Features, Color))
head(df)
```

Dropping the URL, last updated, ad.no columns, and features because they will be no use in doing regression

```
##           Name      Price Model.Year
## 1 Toyota Vitz F 1.0 2017 2385000    2017
## 2 Toyota Corolla GLi Automatic 1.3 VVTi 2019 111000    2019
## 3 Suzuki Alto VXL 2019 1530000    2019
## 4 Suzuki Alto VXR 2019 1650000    2019
## 5 Toyota Corolla XLi VVTi 2010 1435000    2010
## 6 Honda Civic 1.5 RS Turbo 2017 3850000    2017
##           Location Mileage Registered.City Engine.Type
## 1 G- 8, Islamabad Islamabad 9869 Un-Registered 1
## 2 Peshawar KPK 11111 Islamabad 1
## 3 Akora Khattak, Nowshera KPK 17500 Un-Registered 1
## 4 Abdullahpur, Faisalabad Punjab 9600 Lahore 1
## 5 9th Avenue, Islamabad Islamabad 120000 Islamabad 1
## 6 Peshawar Road, Rawalpindi Punjab 22000 Islamabad 1
## Engine.Capacity Transmission Assembly Body.Type
## 1 1000 1 Imported 1
## 2 1300 1 Local 2
```

```
## 3          660          1    Local          1
## 4          660          2    Local          1
## 5         1300          2    Local          2
## 6         1500          1    Local          2
```

```
library(stringr)
df$Name = sub("\\ .*", "", as.character(df$Name))
```

Splitting the name column to only have the value for car brand

```
df <- subset(df, select = -c(Location))
df["Registered"] <- FALSE
df$Registered[df$Registered.City!="Un-Registered"] <- TRUE
df <- subset(df, select = -c(Registered.City))
```

Making registered_city into a True or false column, then dropping registered.city col

```
df["Local"] <- FALSE
df$Local[df$Assembly=="Local"] <- TRUE
df <- subset(df, select = -c(Assembly))
```

Turning imported and local into a true or false column, then dropping the Assembly col

```
df$Registered <- as.factor(df$Registered)
df$Transmission <- as.factor(df$Transmission)
df$Engine.Type <- as.factor(df$Engine.Type)
df$Body.Type <- as.factor(df$Body.Type)
df$Local <- as.factor(df$Local)
df$Name <- as.factor(df$Name)
df[sapply(df, is.integer)] <- lapply(df[sapply(df, is.integer)], as.numeric)
str(df)
```

Making columns into factor and numeric to enhance our model implementation for categorical and integer values

```
## 'data.frame': 46023 obs. of 10 variables:
## $ Name : Factor w/ 31 levels "Adam","Audi",...: 29 29 28 28 29 11 28 22 29 11 ...
## $ Price : num 2385000 111000 1530000 1650000 1435000 ...
## $ Model.Year : num 2017 2019 2019 2019 2010 ...
## $ Mileage : num 9869 11111 17500 9600 120000 ...
## $ Engine.Type : Factor w/ 3 levels "1","2","3": 1 1 1 1 1 1 1 1 2 1 ...
```

```
## $ Engine.Capacity: num 1000 1300 660 660 1300 1500 1000 1000 3000 1800 ...
## $ Transmission : Factor w/ 2 levels "1","2": 1 1 1 2 2 1 2 1 1 1 ...
## $ Body.Type : Factor w/ 6 levels "1","2","3","4",...: 1 2 1 1 2 2 1 1 3 2 ...
## $ Registered : Factor w/ 2 levels "FALSE","TRUE": 1 2 1 2 2 2 2 2 2 2 ...
## $ Local : Factor w/ 2 levels "FALSE","TRUE": 1 2 2 2 2 2 2 1 1 2 ...
```

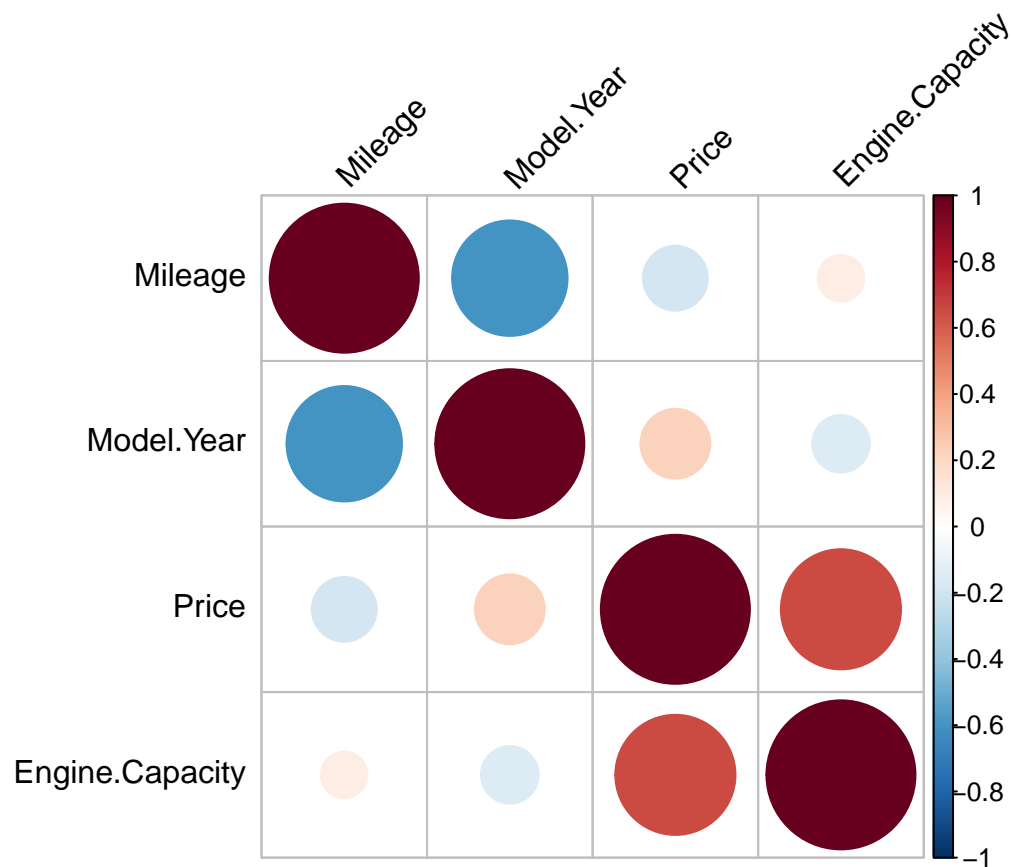
Visual Data exploration

```
library(corrplot)
```

Correlation matrix to identify which numeric columns to use in dataset

```
## corrplot 0.90 loaded
```

```
source("http://www.sthda.com/upload/rquery_cormat.r")
df_numeric <- df[sapply(df, is.numeric)]
rquery.cormat(df_numeric, type="full")
```



```
## $r
##      Mileage Model.Year Price Engine.Capacity
## Mileage      1.000      -0.60 -0.19          0.098
```

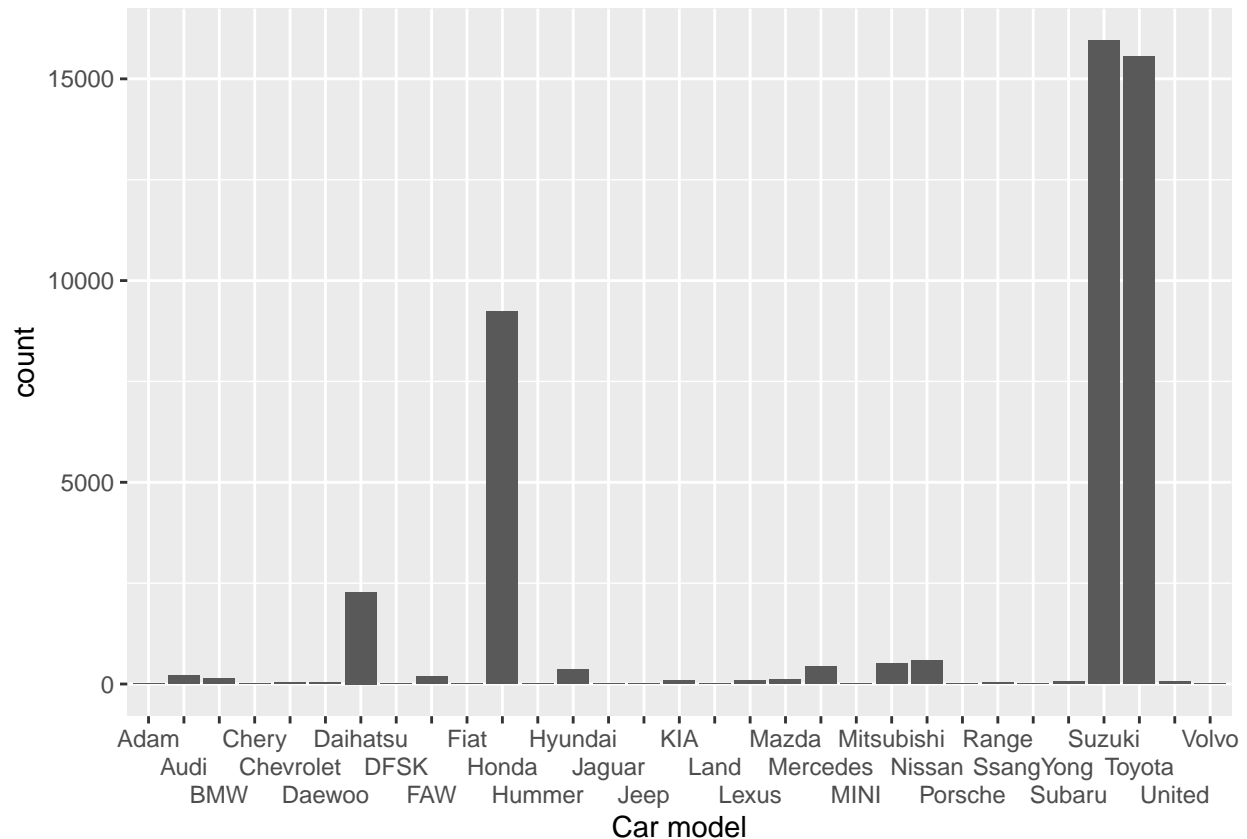
```
## Model.Year      -0.600      1.00  0.22      -0.150
## Price           -0.190      0.22  1.00      0.650
## Engine.Capacity  0.098     -0.15  0.65      1.000
##
## $p
##           Mileage Model.Year Price Engine.Capacity
## Mileage      0.0e+00  0.0e+00   0          7.3e-99
## Model.Year    0.0e+00  0.0e+00   0          9.9e-234
## Price         0.0e+00  0.0e+00   0          0.0e+00
## Engine.Capacity 7.3e-99  9.9e-234   0          0.0e+00
##
## $sym
##           Mileage Model.Year Price Engine.Capacity
## Mileage      1
## Model.Year    .      1
## Price                1
## Engine.Capacity      ,      1
## attr("legend")
## [1] 0 ' ' 0.3 ' ' 0.6 ' ' 0.8 '+' 0.9 '*' 0.95 'B' 1
```

```
library(ggplot2)
```

Barplot to identify which cars are most popular in Pakistan

```
## Warning: package 'ggplot2' was built under R version 4.0.5
```

```
ggplot(df, aes(x = Name)) + geom_bar() + scale_x_discrete(guide = guide_axis(n.dodge=3)) + xlab("Car mo
```



Model Building

Linear Regression

```
set.seed(1234)
spec <- c(train=.6, test=.2, validate=.2)
i <- sample(cut(1:nrow(df),nrow(df)*cumsum(c(0,spec))), labels=names(spec))
train <- df[i=="train",]
test <- df[i=="test",]
vald <- df[i=="validate",]
```

```
lm <- lm(Price~Registered+Transmission+Engine.Type+Body.Type+Local+Mileage+Engine.Capacity+Model.Year, data=train)
summary(lm)
```

I decided to use all the features in dataset because removing the features that were not correlated with price, removed noise from the data and led to lower scores of models.

```
##
## Call:
## lm(formula = Price ~ Registered + Transmission + Engine.Type +
```



```
##      Body.Type + Local + Mileage + Engine.Capacity + Model.Year,
##      data = train)
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -13795014   -498905    -48445     410093    43261392
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -2.455e+08  4.643e+06 -52.881 < 2e-16 ***
## RegisteredTRUE -1.109e+06  4.721e+04 -23.491 < 2e-16 ***
## Transmission2  2.736e+05  3.250e+04  8.418 < 2e-16 ***
## Engine.Type2   -3.963e+06  8.455e+04 -46.864 < 2e-16 ***
## Engine.Type3   -5.505e+05  7.197e+04 -7.649 2.09e-14 ***
## Body.Type2     -5.853e+05  3.163e+04 -18.503 < 2e-16 ***
## Body.Type3      1.863e+06  8.262e+04  22.553 < 2e-16 ***
## Body.Type4      1.995e+04  7.042e+04  0.283  0.777
## Body.Type5      1.617e+05  8.031e+04  2.014  0.044 *
## Body.Type6     -6.306e+05  9.186e+04 -6.865 6.81e-12 ***
## LocalTRUE      -5.313e+05  3.530e+04 -15.051 < 2e-16 ***
## Mileage        -1.968e+00  2.225e-01 -8.846 < 2e-16 ***
## Engine.Capacity 3.310e+03  3.185e+01 103.943 < 2e-16 ***
## Model.Year      1.218e+05  2.300e+03  52.946 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1838000 on 27599 degrees of freedom
## Multiple R-squared:  0.6021, Adjusted R-squared:  0.6019
## F-statistic: 3213 on 13 and 27599 DF, p-value: < 2.2e-16
```

```
# Testing on the data
pred <- predict(lm, newdata = test)
# Computing statistical equation to interpretate our model
print(paste('correlation:', cor(pred, test$Price)))
```

```
## [1] "correlation: 0.771521991335604"
```

```
mse_t <- mean((pred - test$Price)^2)
print(paste("Rmse for test data: ", sqrt(mse_t)))
```

```
## [1] "Rmse for test data: 2051290.16642133"
```

Decision Tree

```
library(tree)
```

```
## Warning: package 'tree' was built under R version 4.0.5
```

```
tree1 <- tree(Price~Name+Registered+Transmission+Engine.Type+Body.Type+Local+Mileage+Engine.Capacity+Model.Year)
summary(tree1)
```

```
##
## Regression tree:
## tree(formula = Price ~ Name + Registered + Transmission + Engine.Type +
##       Body.Type + Local + Mileage + Engine.Capacity + Model.Year,
##       data = train)
## Variables actually used in tree construction:
## [1] "Engine.Capacity" "Name"           "Model.Year"
## Number of terminal nodes: 10
## Residual mean deviance: 1.289e+12 = 3.559e+16 / 27600
## Distribution of residuals:
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -35100000 -444800   -150800      0    371200  41080000
```

```
pred_tree <- predict(tree1, newdata=test)
print(paste('correlation:', cor(pred_tree, test$Price)))
```

Decision Tree performed much better than linear regression because our correlation got alot higher and the RMSE got relatively lower.

```
## [1] "correlation: 0.896548849760948"
```


```
rmse_tree <- sqrt(mean((pred_tree-test$Price)^2))
print(paste('rmse:', rmse_tree))
```

```
## [1] "rmse: 1429424.45115798"
```

Cross validation

```
cv_tree <- cv.tree(tree1)
plot(cv_tree$size, cv_tree$dev, type='b')
```

We will prune the tree to 5 terminal nodes because we want to avoid overfitting by pruning it to



a node with smallest deviance.

Pruning the tree, and then testing.

```
tree_pruned <- prune.tree(tree1, best=5)
pred_pruned <- predict(tree_pruned, newdata=test)
print(paste('correlation:', cor(pred_pruned, test$Price)))
```

In this case, the pruning did not improve results on test data because we got a higher correlation and a lower RMSE for the unpruned Tree.

```
## [1] "correlation: 0.846419959763747"
```

```
rmse_pruned <- sqrt(mean((pred_pruned-test$Price)^2))
print(paste('rmse pruned:', rmse_pruned))
```

```
## [1] "rmse pruned: 1714840.04713052"
```

Support Vector machines

```
library(e1071)
```

```
## Warning: package 'e1071' was built under R version 4.0.5
```

```
svm1 <- svm(Price~Registered+Transmission+Engine.Type+Body.Type+Local+Mileage+Engine.Capacity+Model.Year, data = train, summary(svm1))
```

```
##
## Call:
## svm(formula = Price ~ Registered + Transmission + Engine.Type + Body.Type +
##      Local + Mileage + Engine.Capacity + Model.Year, data = train,
##      kernel = "linear", cost = 10, scale = TRUE)
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: linear
##      cost:   10
##   gamma:    0.07142857
##   epsilon:  0.1
##
## Number of Support Vectors: 10590
```

```
pred <- predict(svm1, newdata=test)
```

```
cor_svm1 <- cor(pred, test$Price)
print(paste('correlation:', cor(pred, test$Price)))
```

SVM got a lower correlation than decision tree and linear regression. The RMSE for SVM was also higher from decision tree and linear regression. I decided not to do hyper parameter tuning for SVM because it took alot of time and was unable to find the optimal paramters, as it gave a warning message “WARNING: reaching max number of iterations”.

```
## [1] "correlation: 0.740897161311206"
```

```
rmse_svm1 <- sqrt(mean((pred - test$Price)^2))
print(paste('rmse:', rmse_svm1))
```

```
## [1] "rmse: 2420531.09540852"
```

Results Analysis

Correlation for these algorthims:

Decision Tree: 0.90

Pruned Decision Tree: 0.85

Linear regression: 0.77

Support vector machine: 0.74

RMSE for these algorithms:

Decision Tree: 1,429,424 (PKR Rupees) or \$6,000

Pruned Decision Tree: 1,714,840 (PKR Rupees) or \$9,600

Linear regression: 2,051,290 (PKR Rupees) or \$12,000

Support vector machine: 2,420,531 (PKR Rupees) or \$12,400

Summary:

Decision tree performed much more efficiently than SVM and linear regression. Our decision tree was off by only about 1.4 million (Rupees) or \$6,000 compared to the other algorithms which were off by more than 2 million (Rupees) or \$12,000. Linear regression works much better when our data is linear, however decision tree works better with more qualitative and factor values and with more complex data. Pruning the decision tree also did not help improve performance of our data. In our case, our data was much more complex because our variables were not much correlated with the price, which is why decision tree performed better. On the other hand, Support vector machine took a lot of time to compile and failed to give efficient results since it is performing dot product of training examples and our data had already been scaled. Furthermore, this script can definitely be useful in the new data because it was able to learn different variations in prices of car based on car's attributes, as we got a correlation of 0.90 and was off by \$6,000.